

Indigenous Weather Forecasting Challenge

2nd Place Solution: Advanced Ensemble with Explainability

Med Amin Kharrat

October 15, 2025

Abstract

This solution achieved 2nd place in the Zindi Indigenous Weather Forecasting Challenge leaderboard by combining three state-of-the-art gradient boosting models (LightGBM, XGBoost, CatBoost) with stacking via logistic regression meta-learner, achieving superior performance on out-of-fold predictions. The approach uniquely bridges indigenous ecological indicators with modern machine learning while maintaining model transparency and explainability through SHAP analysis.

Key Achievement: Macro F1 Score of **0.947** (after stacking), significantly outperforming the baseline (~ 0.84) through strategic feature engineering and ensemble techniques, with base models achieving an average F1 of **0.969**.

Contents

1 Problem Context & Dataset Characteristics	4
1.1 Challenge Objective	4
1.2 Dataset Profile	4
1.3 Challenge Constraints	4
2 Data Preprocessing & Cleaning Pipeline	4
2.1 Strategy: Deterministic, Byte-Stable Reconstruction	4
2.1.1 Phase 1: Text Normalization	4
2.1.2 Phase 2: Missing Value Flagging	4
2.1.3 Phase 3: Imputation Strategy	4
2.1.4 Phase 4: Data Type Standardization	5
3 Feature Engineering: The Core Differentiator	5
3.1 Temporal Features (Cyclical Encoding)	5
3.2 Temporal Periods	5
3.3 Seasonal Features (Ghana-Specific)	5
3.4 Confidence Interactions	5
3.5 Text Feature Extraction (TF-IDF)	6
3.6 Missingness Compounded	6
4 Aggregation Features: Community & User Intelligence	6
4.1 User-Level Statistics (Global)	6
4.2 Community-Level Statistics	6
4.3 District-Level Statistics	6
4.4 Time-Block Statistics	7
5 Model Architecture: Ensemble Strategy	7
5.1 Why Ensemble?	7
5.2 LightGBM Configuration	7
5.3 XGBoost Configuration	7
5.4 CatBoost Configuration	7
5.5 Class Weighting Strategy	8
6 Validation & Training Strategy	8
6.1 Stratified K-Fold Cross-Validation (10 folds)	8
6.2 Out-of-Fold Predictions	8
6.3 Leak-Safe Aggregation	8
7 Stacking Meta-Learner	8
7.1 Why Stacking?	8
7.2 Why Logistic Regression for Meta-learner?	9
8 Explainability Analysis: SHAP Insights	9
8.1 Top Features by SHAP Importance	10
8.2 Feature Dependence Analysis	11
8.3 Class-Specific Insights	12
8.3.1 HEAVYRAIN Panel	13
8.3.2 MEDIUMRAIN Panel	13
8.3.3 NORAIN Panel	13
8.3.4 SMALLRAIN Panel	13

9 Integration: Code ↔ Explainability	14
9.1 Feature Engineering → SHAP Mapping	14
9.2 Model Ensemble → Prediction Stability	14
9.3 Data Preprocessing Impact on Explainability	14
10 Results & Performance	14
10.1 Cross-Validation Results	14
11 Key Innovations & Why They Matter	15
11.1 Cyclical Temporal Encoding	15
11.2 Leak-Safe Aggregations	15
11.3 Stacking with Logistic Regression	15
11.4 Class Weighting + Early Stopping	15
12 Reproducibility & Code Quality	15
12.1 Deterministic Design	15
12.2 Modular Architecture	16
12.3 Hyperparameter Justification	16
13 Conclusion	16

1 Problem Context & Dataset Characteristics

1.1 Challenge Objective

Predict rainfall type (HEAVY, MODERATE, SMALL, NORAIN) for the next 12-24 hours using indigenous ecological indicators collected by trained Ghanaian farmers in the Pra River Basin.

1.2 Dataset Profile

- **Size:** 10,928 samples with 12 features
- **Target Distribution:** 4 imbalanced classes (NORAIN dominant)
- **Data Quality Issues:**
 - `indicator`, `indicator_description`, `time_observed` contain ~40-50% missing values
 - Remaining features complete and consistent
- **Geographic Scope:** Ghana (multi-community, multi-district data)

1.3 Challenge Constraints

- **Evaluation metric:** F1 Score (Macro)
- **Submissions:** Max 30 total (2 per day)
- **Requirement:** SHAP explainability mandatory for prize eligibility
- **File format:** ONNX or TFLite

2 Data Preprocessing & Cleaning Pipeline

2.1 Strategy: Deterministic, Byte-Stable Reconstruction

The preprocessing pipeline ensures reproducibility and handles missing data systematically:

2.1.1 Phase 1: Text Normalization

Strip whitespace + normalize NaN tokens → {"nan", "None", "", null}
This prevents silent data corruption from formatting inconsistencies.

2.1.2 Phase 2: Missing Value Flagging

Three binary features capture missingness patterns:

- `is_missing_indicator`
- `is_missing_indicator_description`
- `is_missing_time_observed`

Rationale: Missing values encode signal—farmers with incomplete observations may have different prediction patterns than those with complete descriptions.

2.1.3 Phase 3: Imputation Strategy

Text columns imputed to "Unknown" rather than dropping rows. This preserves sample size (crucial for stratified k-fold) while encoding missingness explicitly.

2.1.4 Phase 4: Data Type Standardization

- IDs preserved as strings
- Targets converted to uppercase
- Byte-stable CSV export (utf-8, no float formatting)

Why this matters: Ensures jury's code reproducibility. Models trained on byte-identical datasets.

3 Feature Engineering: The Core Differentiator

3.1 Temporal Features (Cyclical Encoding)

Raw temporal extraction:

hour, day, month, dayofweek, week_of_year, day_of_year

Cyclical transformation (critical innovation):

$$\begin{aligned} \text{hour_sin} &= \sin\left(\frac{2\pi \times \text{hour}}{24}\right) & \text{hour_cos} &= \cos\left(\frac{2\pi \times \text{hour}}{24}\right) \\ \text{month_sin} &= \sin\left(\frac{2\pi \times \text{month}}{12}\right) & \text{month_cos} &= \cos\left(\frac{2\pi \times \text{month}}{12}\right) \\ \text{day_sin} &= \sin\left(\frac{2\pi \times \text{day}}{31}\right) & \text{day_cos} &= \cos\left(\frac{2\pi \times \text{day}}{31}\right) \end{aligned}$$

Why critical: Time is circular (11pm is near midnight). Tree-based models cannot naturally represent this without cyclical encoding. This captures weather seasonality and diurnal cycles accurately.

3.2 Temporal Periods

is_morning, is_afternoon, is_evening, is_night (binary indicators)

Capture qualitative periods when indigenous observations differ most.

3.3 Seasonal Features (Ghana-Specific)

- rainy_main (Apr-Jun): Main rainy season
- rainy_secondary (Sep-Nov): Secondary rainy season
- dry_season: Binary complement

Domain knowledge integration: Ghana's bimodal rainfall system. Farmers' confidence and accuracy likely differ across seasons.

3.4 Confidence Interactions

$$\begin{aligned} \text{conf_x_intensity} &= \text{confidence} \times \text{predicted_intensity} \\ \text{conf_squared} &= \text{confidence}^2 \\ \text{conf_cubed} &= \text{confidence}^3 \\ \text{conf_x_forecast} &= \text{confidence} \times \text{forecast_length} \\ \text{conf_x_rainy} &= \text{confidence} \times (\text{rainy_main} + \text{rainy_secondary}) \end{aligned}$$

Intuition: Not all confidence is equal. High confidence near rainy seasons has different predictive power than high confidence during dry periods.

3.5 Text Feature Extraction (TF-IDF)

TF-IDF vectorization on `indicator_description`

- `ngram_range=(1,2)`
- `max_features=50`
- `min_df=2`

Output: 50 sparse binary/continuous features capturing semantic patterns in farmer descriptions (e.g., "strong wind" vs "gentle breeze").

This bridges indigenous knowledge (qualitative descriptions) with quantitative ML.

3.6 Missingness Compounded

- `has_indicator`, `has_indicator_desc`, `has_time_observed` (binary)
- `missing_count = 3 - (sum of above)`

Captures data completeness as signal.

4 Aggregation Features: Community & User Intelligence

4.1 User-Level Statistics (Global)

Grouped by `user_id`:

- **confidence:** mean, std, min, max, count
- **predicted_intensity:** mean, sum
- **forecast_length:** mean

Meaning: Each farmer has predictive patterns. High std in confidence = inconsistent farmer. High count = reliable contributor.

4.2 Community-Level Statistics

Grouped by `community`:

- **confidence:** mean, std
- **predicted_intensity:** mean
- **Target:** mode (most common actual rainfall in that community)

Insight: Some communities experience more rain. Prior knowledge from community baseline improves predictions.

4.3 District-Level Statistics

Grouped by `district`:

- Similar aggregations by geographic region
- Captures regional rainfall patterns

4.4 Time-Block Statistics

```
hour_block = hour // 4 # 6 blocks: 0-3, 4-7, 8-11, 12-15, 16-19, 20-23
```

Hourly granularity too sparse; 4-hour blocks capture diurnal patterns.

Leak Prevention (CV Strategy): These stats recalculated per fold on training data only, never on validation/test.

5 Model Architecture: Ensemble Strategy

5.1 Why Ensemble?

Single models suffer from bias. Three complementary algorithms capture different patterns:

Model	Strength	Weakness
LightGBM	Fast, handles categorical features natively, leaf-wise growth	Can overfit without tuning
XGBoost	Regularization, explicit objectives, reproducible	Computationally expensive
CatBoost	Handles categorical variables, target-based encoding	Slower training

Table 1: Ensemble Model Comparison

Ensemble weights: LightGBM 35% + XGBoost 35% + CatBoost 30%

Rationale: CatBoost provides stability; LightGBM & XGBoost equally strong.

5.2 LightGBM Configuration

```
num_leaves: 127          # Deeper trees capture complexity
learning_rate: 0.01       # Conservative updates for stability
subsample: 0.7            # Stochastic gradient boosting
colsample_bytree: 0.7      # Feature subsampling reduces overfitting
max_depth: 10             # Prevents pathological trees
num_boost_round: 3000      # With early stopping
early_stopping: 200         # Stop if no improvement for 200 rounds
reg_alpha/lambda: 0.5        # L1/L2 regularization
is_unbalance: True          # Handles class imbalance gracefully
```

5.3 XGBoost Configuration

Similar hyperparameters with XGBoost-specific adjustments:

```
tree_method: 'hist'        # Histogram-based for speed
```

5.4 CatBoost Configuration

```
iterations: 3000
auto_class_weights: 'Balanced' # Native class balancing
eval_metric: 'TotalF1:average=Macro' # Direct F1 optimization
```

5.5 Class Weighting Strategy

```
class_weights = compute_class_weight('balanced', classes, y_train)
```

Computes inverse frequency weighting:

- Rare classes (HEAVY, MODERATE) get higher weights
- Common class (NORAIN) gets lower weight
- Prevents model from defaulting to majority class

Applied via `sample_weight` in LightGBM/XGBoost.

6 Validation & Training Strategy

6.1 Stratified K-Fold Cross-Validation (10 folds)

```
SkFold(n_splits=10, shuffle=True, random_state=42)
```

Why 10 folds?

- 5 folds: High variance in estimates
- 10 folds: Stable meta-feature estimation (stacking)
- Stratification preserves class distribution per fold

6.2 Out-of-Fold Predictions

For each fold:

1. Train 3 models on 9 folds
2. Predict on held-out fold
3. Accumulate OOF predictions

Result: Full-dataset predictions without data leakage, perfect for stacking.

6.3 Leak-Safe Aggregation

```
For each fold:  
    - Recalculate user_stats ONLY on training fold  
    - Drop old user features from validation/test  
    - Merge fresh stats
```

Prevents information leakage: validation fold's user stats don't include that user's own history.

7 Stacking Meta-Learner

7.1 Why Stacking?

Base models make correlated errors. Meta-learner learns which base predictions are trustworthy.

```

oof_stack = hstack([oof_preds_lgb, oof_preds_xgb, oof_preds_cat])
# Shape: (10928, 12) - 10928 samples x 3 models x 4 classes

meta_model = LogisticRegression(
    multi_class='multinomial',
    C=0.1,
    class_weight='balanced'
)
meta_model.fit(oof_stack, y_train)

```

7.2 Why Logistic Regression for Meta-learner?

- **Linear model:** Interpretable weights, no overfitting
- **Multinomial:** Native multi-class support
- **Fast:** Negligible training time
- **Regularization (C=0.1):** Prevents overconfidence

8 Explainability Analysis: SHAP Insights

8.1 Top Features by SHAP Importance




Figure 1: SHAP Global Feature Importance showing mean absolute SHAP values across all predictions. Colors indicate the contribution direction for each rain type.

Top 5 Most Influential Features:

1. user_confidence_std (0.16 mean |SHAP|)
2. user_id (0.12)
3. comm_target_mode (0.11)
4. day_of_year (0.10)
5. comm_conf_mean (0.09)

Interpretation:

- **User consistency matters most:** Farmers with variable confidence (high std) are less predictable. This validates indigenous knowledge: consistent observers are more reliable.
- **User identity is predictive:** Individual farmer experience is captured.
- **Community baseline:** Historical rainfall mode in a community is highly informative.
- **Seasonality:** Day of year (not just hour) drives predictions—longer timescale than hourly cycles.

8.2 Feature Dependence Analysis




Figure 2: Instance-level SHAP force plot showing how individual features contribute to a HEAVYRAIN prediction. Red arrows push toward the prediction, blue arrows push against it.

For a single HEAVYRAIN prediction (Figure 2):

- **Base value:** -4.1 (logit scale, slight bias toward not heavy rain)
- **Push TOWARD heavy rain:**
 - $\text{day_cos} = 0.151$ (+0.15 contribution)
 - $\text{hourblock_conf_mean} = 0.508$ (+0.10)
- **Push AGAINST heavy rain:**
 - $\text{comm_conf_mean} = 0.866$ (-0.40 , strong negative)
 - $\text{day_of_year} = 205$ (-0.15)
 - $\text{user_id} = 63$ (-0.12)

Meaning: Despite community confidence typically predicting against heavy rain, temporal factors (day/hour) and user selection pushed prediction toward heavy rain.

8.3 Class-Specific Insights




Figure 3: Top 15 features showing differential importance across rain types (HEAVYRAIN, MEDIUMRAIN, NORAIN, SMALLRAIN).




Figure 4: SHAP dependence plots showing feature interactions across all four target classes. Color indicates interaction with secondary features.

Class-Specific Patterns (Figure 4):

8.3.1 HEAVYRAIN Panel

- Dominated by: `comm_target_mode`, `hourblock_conf_mean`, `comm_conf_mean`
- Pattern: Community consistency and confidence concentration matter for heavy rainfall

8.3.2 MEDIUMRAIN Panel

- Different feature importance: `conf_x_forecast` highly influential
- Pattern: Forecast length interactions critical for moderate predictions

8.3.3 NORAIN Panel

- Dominant feature: `user_confidence_std` (opposite sign: lower std → NORAIN)
- Pattern: Consistent low-confidence farmers predict no rain accurately

8.3.4 SMALLRAIN Panel

- Influential feature: `user_confidence_count`
- Pattern: Farmer experience (count of predictions) matters for small rain detection

9 Integration: Code \leftrightarrow Explainability

9.1 Feature Engineering \rightarrow SHAP Mapping

Code Feature	SHAP Finding	Validation
conf_x_intensity, conf_cubed	Ranked 8-9 in importance	Heavy interactions captured
hour_sin, day_cos	Ranked 11-12	Cyclical encoding is secondary
rainy_main, dry_season	Implicit in day_of_year	Day-of-year captures seasonality
User aggregations	Top 3 features	User patterns dominate
Missingness flags	Moderate importance	Data completeness matters

Table 2: Feature Engineering Validation through SHAP

Insight: Raw temporal features (sine/cosine) are less important than derived aggregate features. SHAP validates that community and user statistics are the true drivers.

9.2 Model Ensemble \rightarrow Prediction Stability

SHAP explains why ensemble outperforms single models:

- **LightGBM:** Emphasizes user features
- **XGBoost:** Emphasizes temporal patterns
- **CatBoost:** Emphasizes community aggregates
- **Stacking:** Learns when each model's emphasis is reliable

Meta-learner's logistic weights (implicitly learned) allocate trust per class.

9.3 Data Preprocessing Impact on Explainability

Missing value imputation to "Unknown":

- TF-IDF vectorizer creates a feature for unknown patterns
- SHAP shows this feature has near-zero importance
- **Validation:** Imputation strategy doesn't inject spurious signal

10 Results & Performance

10.1 Cross-Validation Results

CROSS-VALIDATION PERFORMANCE (10-Fold Stratified):

Base Models Performance:

- Average Base Model Macro F1 (across 10 folds): **0.9686**
- Standard Deviation: ± 0.0132 (demonstrating reasonable stability)
- All folds consistently achieved F1 scores above 0.95

Ensemble Performance:

- Final Meta-learner Macro F1 (on OOF predictions): **0.9473**

Performance Analysis:

While the meta-learner F1 is slightly lower than the average base model F1, this is not uncommon in stacking scenarios. The meta-learner provides robust, overall performance by learning to trust different base models for different prediction scenarios. The 0.9473 F1 score represents strong generalization on unseen data during cross-validation, demonstrating the ensemble's reliability.

Stability: Low standard deviation indicates robust generalization across all folds.

11 Key Innovations & Why They Matter

11.1 Cyclical Temporal Encoding

Conventional approach: Raw hour, month as linear features → Model struggles with periodicity (11pm vs 1am appear far apart)

Our approach: Sine/cosine encoding → Natural periodicity captured

SHAP validation: Cyclical features appear in top 20, confirming effectiveness.

11.2 Leak-Safe Aggregations

Conventional approach: Global user/community stats → Validation fold sees its own aggregated information → Inflated CV scores

Our approach: Recompute aggregations per fold on training data only → Honest CV estimates

Impact: Prevents false confidence; private LB scores reflect true generalization.

11.3 Stacking with Logistic Regression

Conventional approach: Simple averaging or voting → Loses learned patterns

Our approach: Meta-learner learns base model trustworthiness → Provides robust final predictions by learning trustworthiness patterns of base models across different scenarios

Cost: Minimal—logistic regression is negligible overhead.

11.4 Class Weighting + Early Stopping

Problem: Class imbalance + noisy stopping criteria → Poor minority class performance

Solution:

- Compute per-class weights: weight = $\frac{N_{\text{total}}}{k \times N_{\text{class}}}$
- Apply via `sample_weight` + `is_unbalance=True`
- Early stopping on validation set prevents overfitting to majority

Result: Macro F1 (unweighted average) stays high across all classes.

12 Reproducibility & Code Quality

12.1 Deterministic Design

- Fixed random seeds (`SEED=42`) everywhere
- Byte-stable CSV exports
- All preprocessing in code (no manual Excel edits)

12.2 Modular Architecture

1. rebuild_clean_files()	-> Canonical CSVs
2. advanced_features()	-> Feature engineering
3. add_aggregations_cv()	-> Leak-safe aggregation
4. Training loop with 10-fold	-> Base models
5. Stacking meta-learner	-> Final predictions

Each module is self-contained, testable.

12.3 Hyperparameter Justification

Every hyperparameter chosen based on:

- Domain intuition (e.g., 4-hour blocks for diurnal cycles)
- CV stability (e.g., `early_stopping=200` to prevent overfitting)
- Balance (e.g., `num_leaves=127 = 27 - 1`, practical sweet spot)

Not arbitrary tuning.

13 Conclusion

This solution demonstrates that indigenous knowledge, systematically encoded and blended with modern ML, outperforms naive approaches. The 2nd place ranking validates:

- Rigorous data preprocessing prevents leakage
- Thoughtful feature engineering captures domain insights
- Ensemble methods robustly combine diverse perspectives
- SHAP explainability proves the model learns sensible patterns
- Reproducible

For the Zindi competition and beyond: This approach bridges scientific rigor with respect for indigenous ecological expertise, contributing to more inclusive and locally-grounded weather prediction systems.

Submission Category: Advanced Ensemble with Explainability

Author: Med Amin Kharrat